CS 589 Fall 2020

Information retrieval infrastructure

Instructor: Susan Liu
TA: Huihui Liu

Stevens Institute of Technology

Information Retrieval Infrastructure

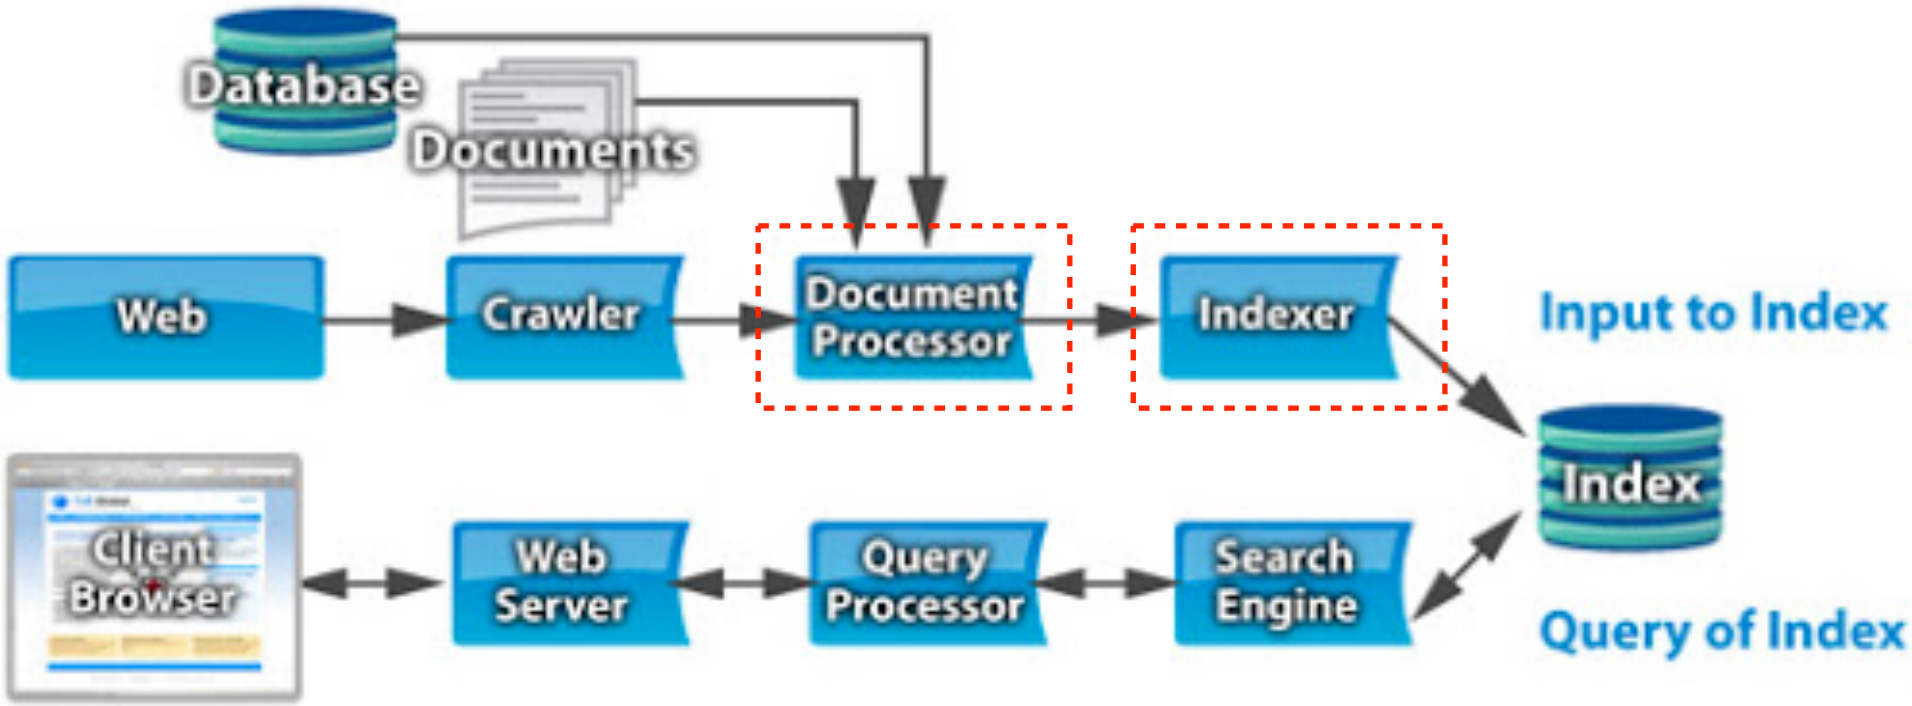

Inverted index

- In Lecture 2, we learned retrieval models
  - Compute score(q, d)
  - Select the d that maximizes score(q, d)

- In an industry scale search engine, there could be trillions of q's and billions of d's
  - For each query, search time complexity = O(|D|)
  - Solution for faster retrieval: inverted index

Inverted index

1: Winter is coming.
2: Ours is the fury.
3: The choice is yours.

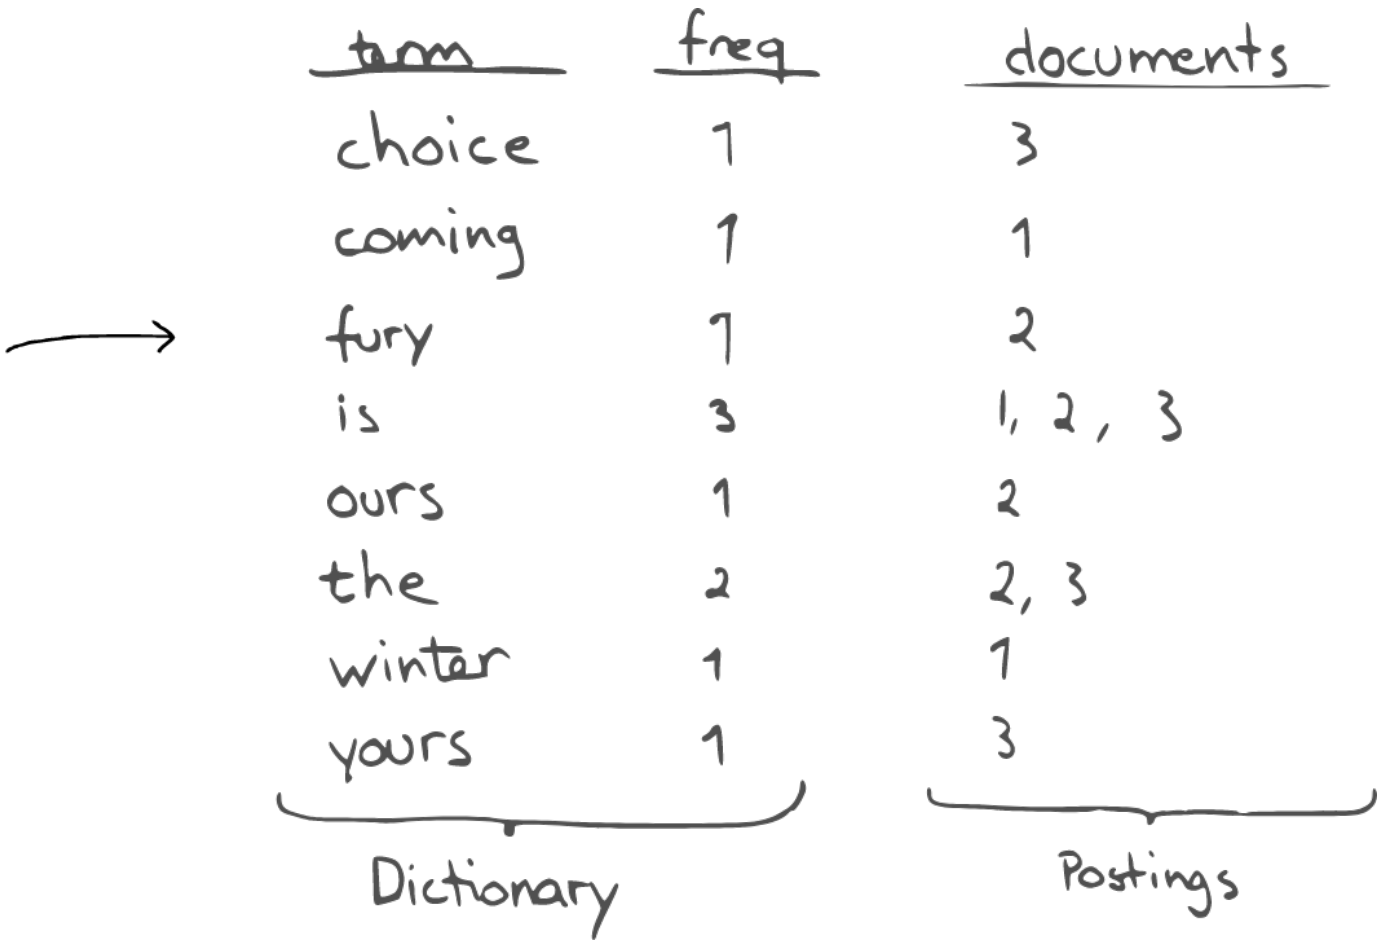

time complexity: O(#unique words in q x avg_len(postings lists))

$$\ll |D|$$

Problems with inverted indexing

- Data processing
  - Choosing the unit for indexing
  - Determining the vocabulary
- Constructing/speeding up inverted index
  - Skipping index
  - Prefix indexing
  - Indexing with blocks
  - MapReduce
- Index compression
- Other issues
  - Indexing position
  - Spelling correction

Choosing the correct unit for indexing

- Documents often consists of sub documents
  - e.g., email contains multiple attached documents

- Trade-off on the unit size
  - Smaller units: missing important passages
  - Larger unit: gets spurious matches, e.g., ….**text** messages….gold **mining**…

Determining the vocabulary

- Tokenization

Input: Friends, Romans, Countrymen, lend me your ears;
Output: | Friends | Romans | Countrymen | lend | me | your | ears |

  - o'neil, aren't, C#

- Dropping stop words
  - Stop words are common terms
  - **Web search engines generally do not use stop words!**

A	It	These
About	Its	They
Again	Itself	This
All	Just	Those
Almost	km	Thus
Also	Made	To
Although	Mainly	Upon
Always	Make	Use
An	May	Used
And	mg	Using
Another	Might	Various
Any	ml	Very
Are	mm	Was
As	Most	We
At	Mostly	Were

Determining the vocabulary

- Normalization
  - Abbrev: USA vs. United states of America
  - Case:
    - Cat -> cat
    - **SAT** -> **sat**

- Stemming/lemmatization
  - singing -> sing, cars -> car, **sat -> sit**
  - porter stemmer, snowball stemmer

Speeding up: skipping lists

- Finding the intersection of two post listings
  - Without skip: O(m + n)

```
i, j = 0, 0
while i < m and j < n:
    if arr1[i] < arr2[j]:
        i += 1
    elif arr2[j] < arr1[i]:
        j+= 1
    else:
        print(arr2[j])
        j += 1
        i += 1
```

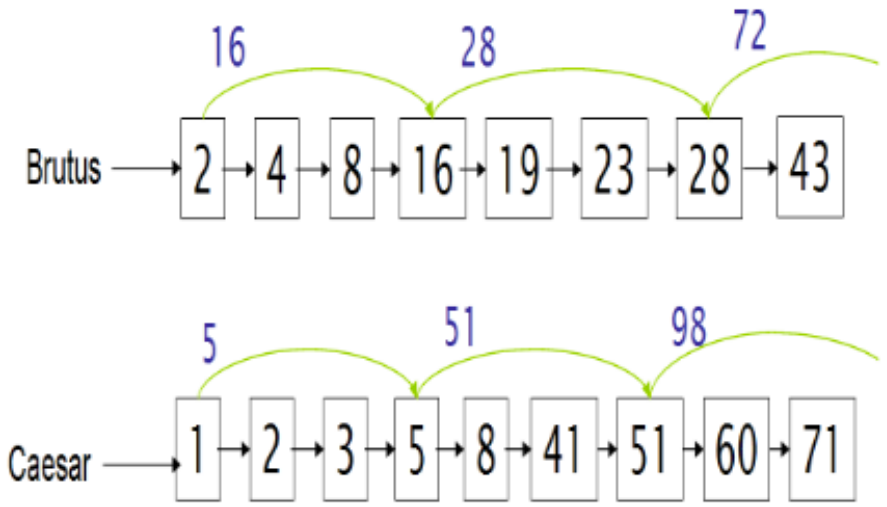

Speeding up: prefix indexing

- Speeding up the indexing using prefix tree
  - time complexity: O(**#unique words in q** x avg_len(postings lists))

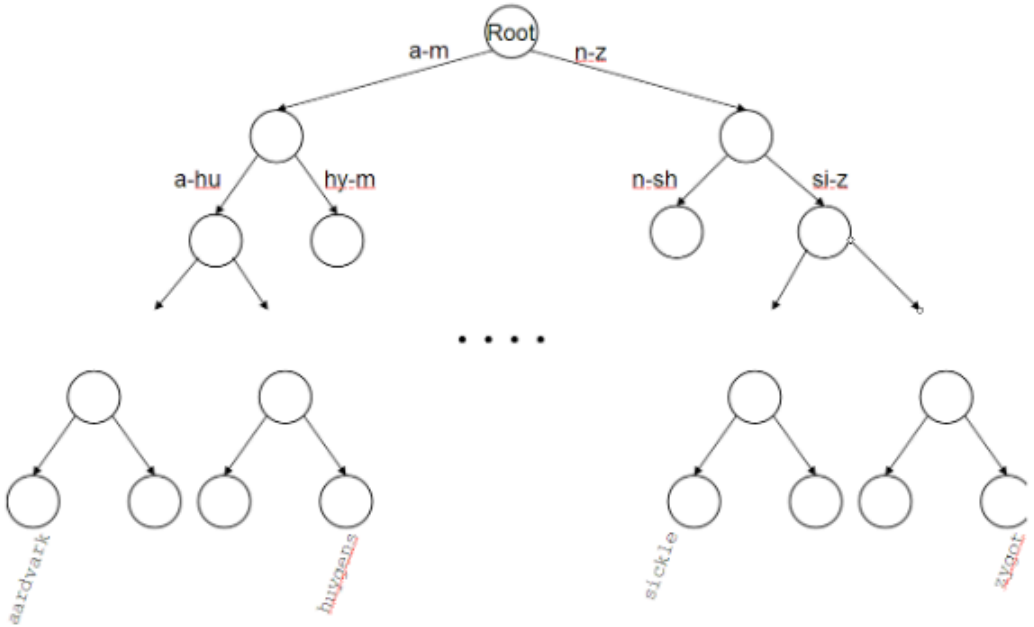

Constructing inverted index: hardware basics

- Decisions on an IR system largely depends on the hardware which the system runs on

- **Chunks**:
  - Splitting data into more chunks takes more *seek time*

- **Blocks**
  - *Accessing data in memory >> accessing data on disk*
  - *Constructing* inverted index using blocks
  - Typical IR system: GBs of memory, disk space orders of magnitude larger

Block sort-based indexing

- Indexing large corpus
  - Reuters-RCV1: 2.5GB = 2.5 x 10^9Bytes, 1 billion
  - Today's text corpus contains petabytes of data: 10^15Bytes
  - Memory << size of corpus

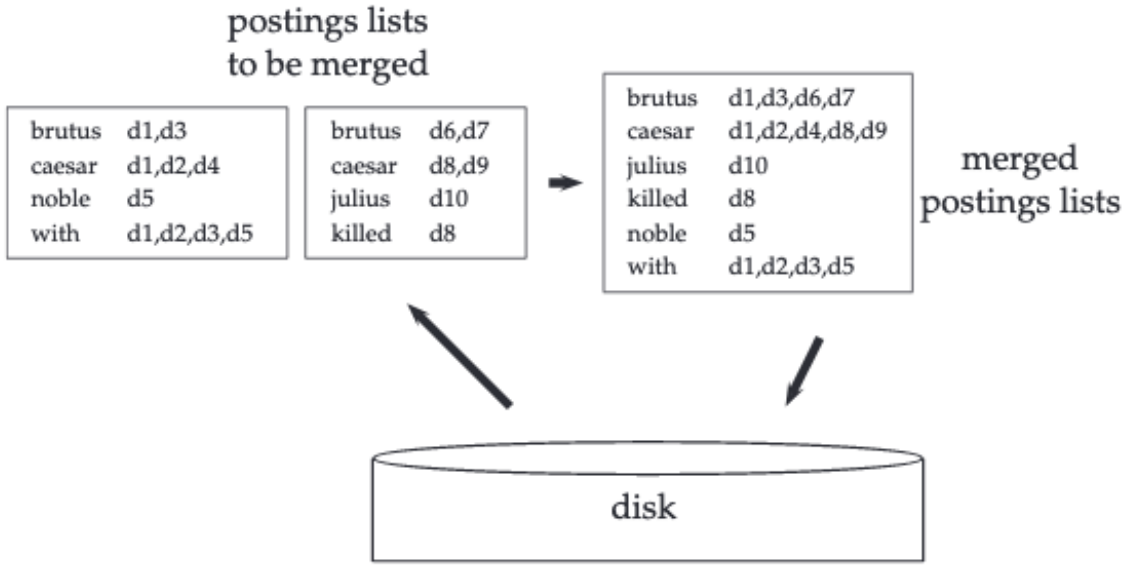

- Index each block using memory
- Write each blocks' index into disk
- Merge all inverted indices

Block sort-based indexing

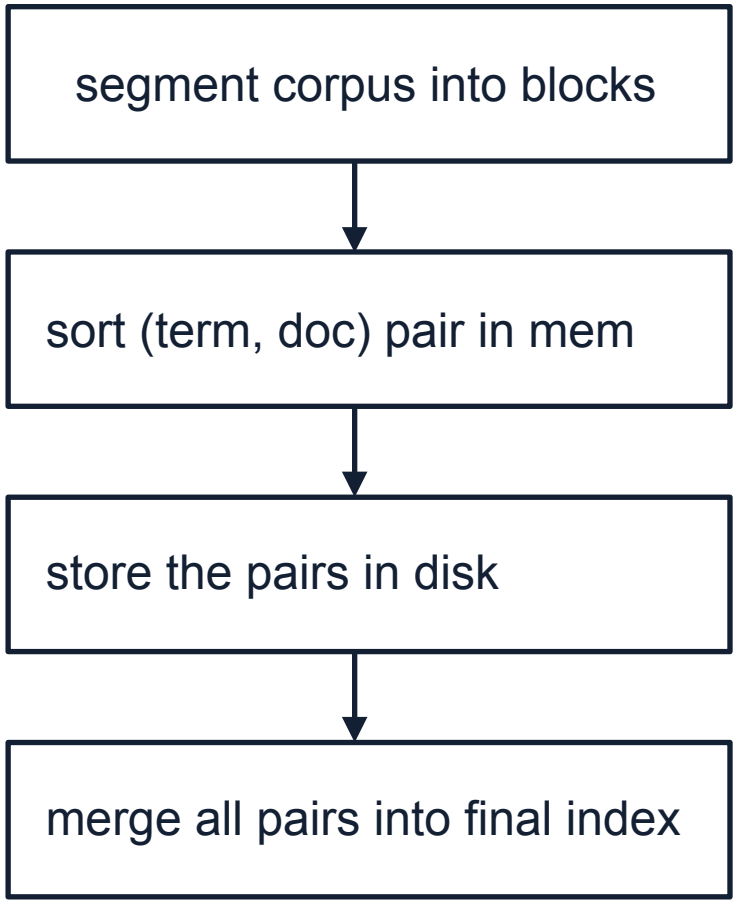


```
BSBINDEXCONSTRUCTION()
1  n ← 0
2  while (all documents have not been processed)
3  do n ← n + 1
4     block ← PARSENEXTBLOCK()
5     BSBI-INVERT(block)
6     WRITEBLOCKTODISK(block, f_n)
7  MERGEBLOCKS(f_1, ..., f_n; f_merged)
```

Single-pass in-memory indexing

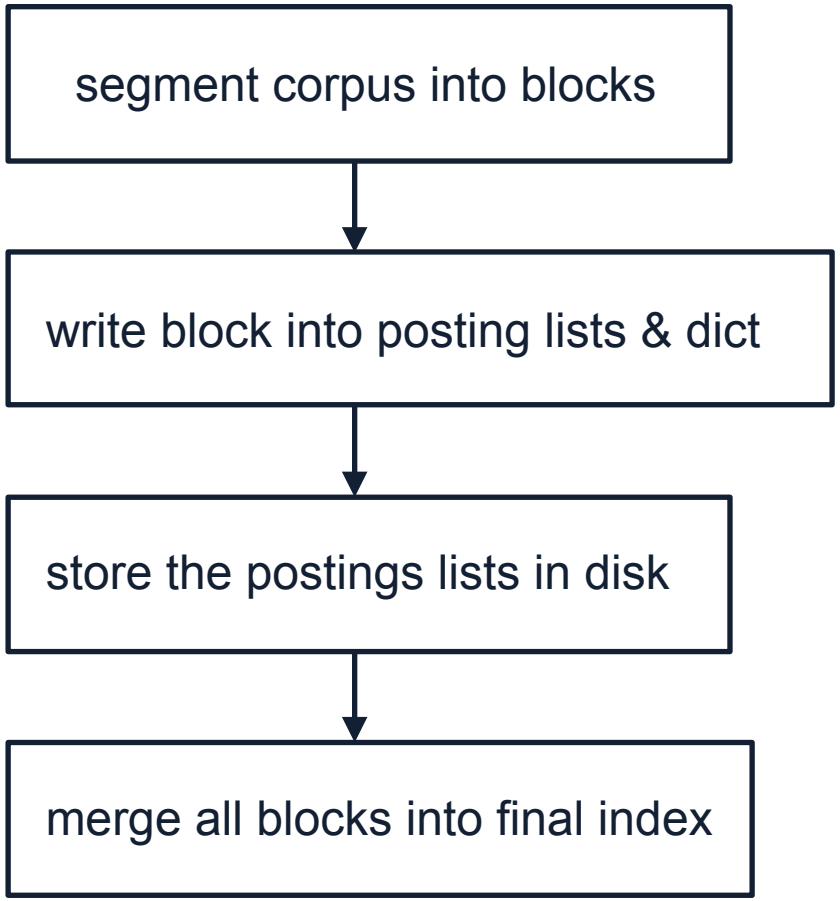

- Handling posting lists directly
- Eliminating the expensive sorting in BSBI
- Leveraging compression

Handling web scale indexing

- Web-scale indexing must use clusters of servers
  - Google had 1 million servers in 2011

- Fault tolerance of a massive data center
  - If a non-fault tolerance system has 1000 nodes, each has 99.9% uptime, then 63% of the time one or more servers is down

- Solution
  - Maintain a “master” server
  - Break indexing into parallel tasks
  - Assign each task to an idle machine

Map-reduce

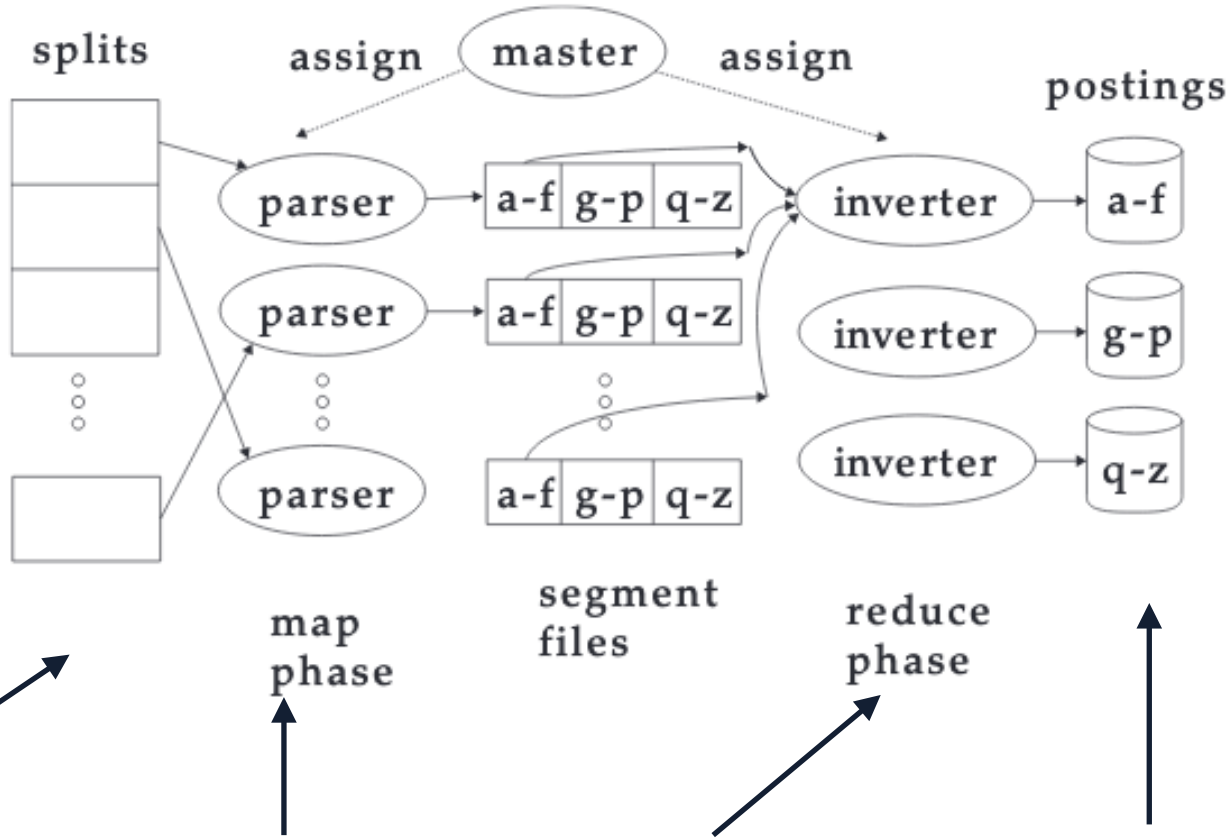

master assigns split to idle machine

parser emits (term,doc) pair

merge partitions in inverter

complete the index

Examples of map-reduce

map: d_2 : C died. d_1 : C came, C c'ed.

(⟨C, d_2⟩, ⟨died,d_2⟩, ⟨C,d_1⟩, ⟨came,d_1⟩, ⟨C,d_1⟩, ⟨c'ed,d_1⟩)

reduce: (⟨C,(d_2,d_1,d_1)⟩,⟨died,(d_2)⟩,⟨came,(d_1)⟩,⟨c'ed,(d_1)⟩)

(⟨C,(d_1:2,d_2:1)⟩,⟨died,(d_2:1)⟩, ⟨came,(d_1:1)⟩, ⟨c'ed,(d_1:1)⟩)

MapReduce: Industry practice

- Term partition vs. document partition
  - Term-partitioned: one machine handles a subrange of terms
  - Document-partitioned: one machine handles a subrange of documents

- Most industry search engine use document-partitioned index
  - Better load balancing (**why?**)

MapReduce: Simplified Data Processing on Large Clusters

Jeffrey Dean and Sanjay Ghemawat

jeff@google.com, sanjay@google.com

Google, Inc.

Dynamic indexing

- Document collection are updated all the time

- How to handle dynamic indexing?
  - Maintain one-big index
  - New document goes to auxiliary “smaller” index
  - Search across both, merge results
  - Periodically merge the two indices
  - Deletion: maintain bitvectors of deleted documents

Real time search of Twitter

- Requires high real time search
  - Low latency, high throughput query evaluation
  - High ingestion rate and immediate data availability
  - Concurrent reads and writes of the index

- Solution: using segments
  - Each segment consists of 2^32 tweets (in memory)
  - New posts are appended to the posting lists
  - Only one segment can be written to at each time

Index compression

- Why compression?
  - Using less disk space
  - Compressing dictionary
    - Allowing the dictionary to be stored in memory
  - Compressing posting files
    - Reducing disk space

- Zipf's law
  - The ith most frequent term has frequency proportional to 1/i

Dictionary compression

- Most of the space in the table is wasted
  - Most words are no 20 bytes
  - Table storage = 28N

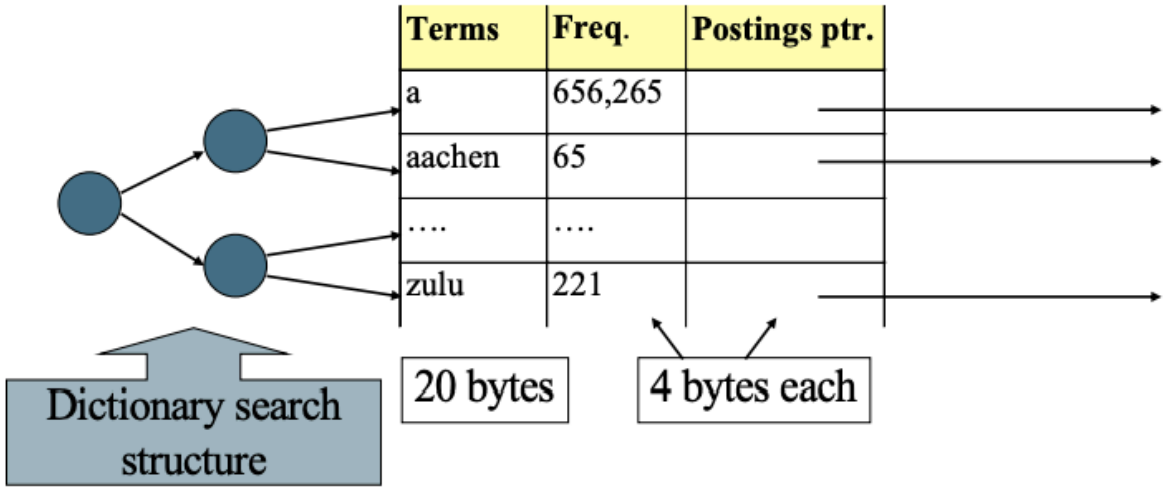

Dictionary-as-a-string

- Table storage = 11N

- How to further improve the storage space?
  - Instead of storing absolute term pointers, store the gaps

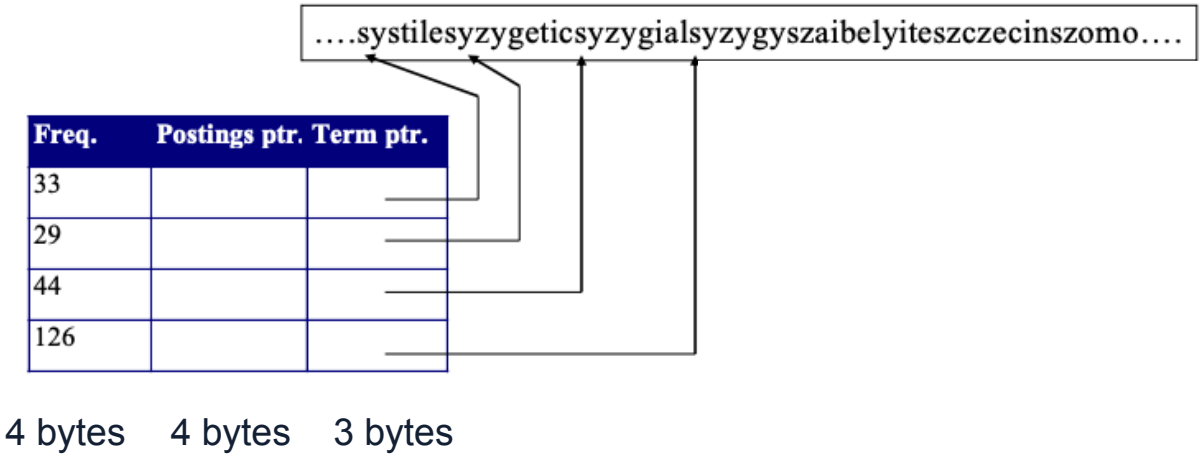

Dictionary-as-a-string

- Table storage = 8N + 3N * (7/12) = 9.75N < 11N
- Trade-off between skipping more vs. skipping less

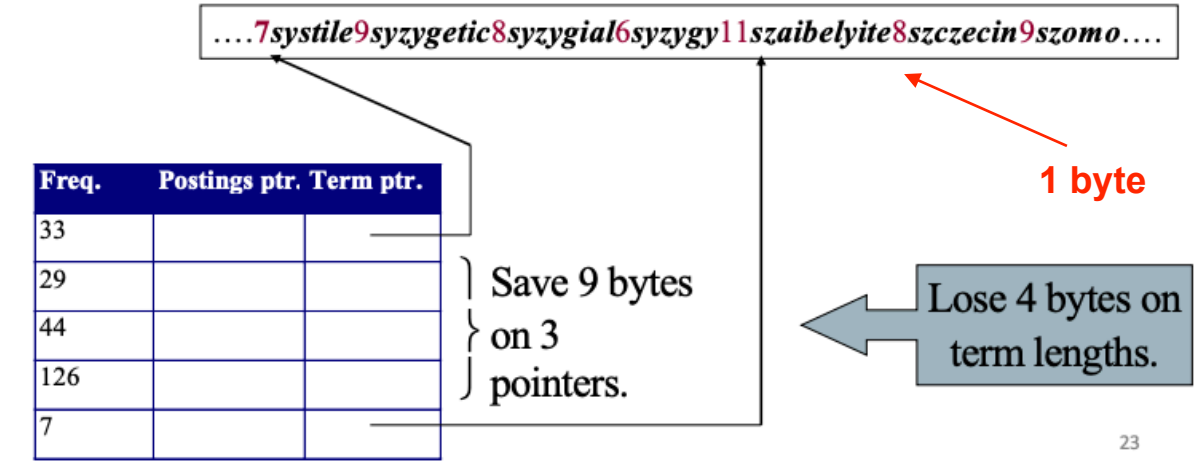

4 bytes 4 bytes 3 bytes

Postings compression

- Observations of posting files
  - Instead of storing docID, store gaps
  - Brutus: 2,4,8,3,4,5,15
  - Binary seq: 10,100,1000,11,100,101,1

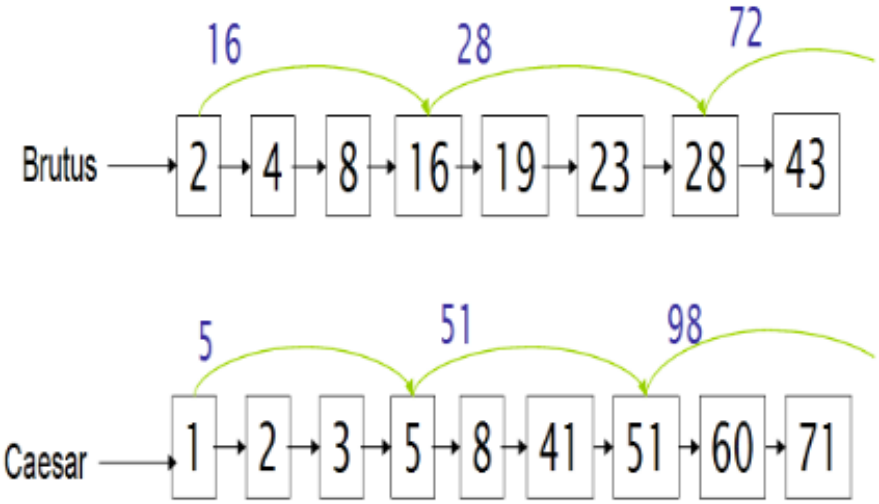

- Prefix encoding
  - Binary encoding such that the sequence can be uniquely decoded
  - e.g., Huffman encoding
  - Unary encoding: {2:110,4:11110, …}
  - A uniquely decodable seq: 1101111101111111101110…

Postings compression

- Problem with unary encoding
  - Too long!

- Gamma code of 13: 1110,101
  - Length: 1110
  - Offset: 13 → 1101 → 101

number	length	offset	γ-code
0			none
1	0		0
2	10	0	10,0
3	10	1	10,1
4	110	00	110,00
9	1110	001	1110,001
13	1110	101	1110,101
24	11110	1000	11110,1000
511	111111110	11111111	111111110,11111111
1025	11111111110	0000000001	11111111110,0000000001

- We can prove gamma code is uniquely decodable

- Gamma code compression rate: 11.7%

Indexing Position

- Indexing the position of word within the document
- Intersection algorithm finds where the two terms appear between within *k* words

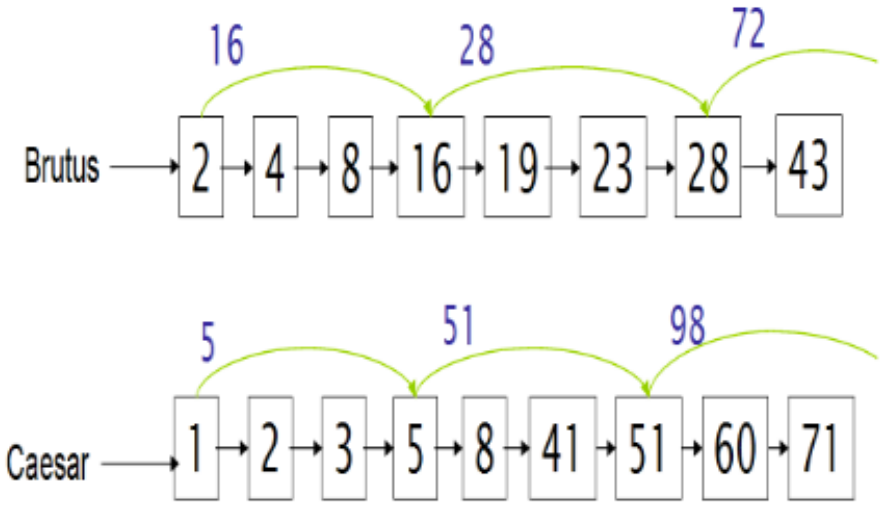

Brutus: 2:<0>, 4: <429,433>, 8: <150>, …

Ceasar: 1:<10>, 2: <29>, 8: <17, 250>, …

Spelling correction

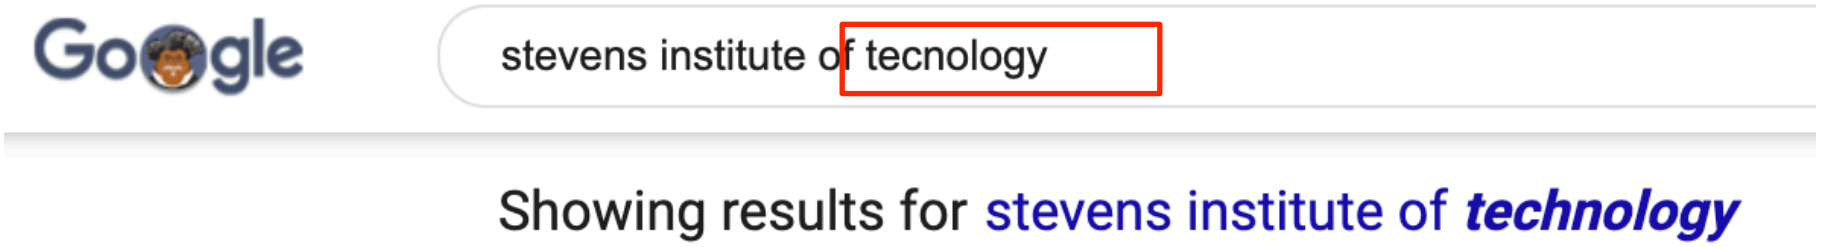

- Edit distance
- k-gram index for spelling correction
- context sensitive spelling correction

Edit distance

- Dynamic programming: O(|s1| x |s2|)

```
EDITDISTANCE(s1, s2)
 1  int m[i, j] = 0
 2  for i ← 1 to |s1|
 3  do m[i, 0] = i
 4  for j ← 1 to |s2|
 5  do m[0, j] = j
 6  for i ← 1 to |s1|
 7  do for j ← 1 to |s2|
 8       do m[i, j] = min{m[i − 1, j − 1] + if (s1[i] = s2[j]) then 0 else
 9                        m[i − 1, j] + 1,
10                        m[i, j − 1] + 1}
11  return m[|s1|, |s2|]
```

			f		a		s		t	
		0	**1**	**1**	**2**	**2**	**3**	**3**	**4**	**4**
c	**1**	**1**	*1*	*2*	**2**	3	**3**	4	**4**	5
	1	**1**	*2*	*1*	**2**	**2**	**3**	**3**	**4**	**4**
a	**2**	**2**	**2**	**2**	*1*	*3*	*3*	*4*	4	5
	2	**2**	3	**2**	*3*	*1*	*2*	*2*	**3**	**3**
t	**3**	**3**	**3**	**3**	3	**2**	*2*	*3*	*2*	*4*
	3	**3**	4	**3**	4	**2**	*3*	*2*	*3*	*2*
s	**4**	**4**	**4**	**4**	4	**3**	**2**	3	*3*	*3*
	4	**4**	5	**4**	5	**3**	4	**2**	*3*	*3*

Levenshtein distance

k-gram indexes for spelling correction

- Running DP on all pairs of words is time consuming
- Leveraging k-gram index to speed up spelling correction
  - boardroom vs. bord

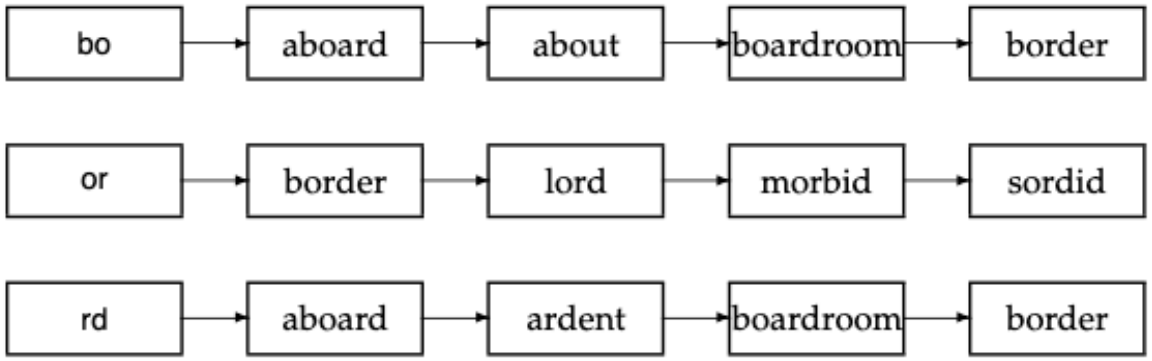

boarder:3

boardroom: 2

aboard: 2

ardent:1

…

Context sensitive spelling correction

- How to correct “flew form healthrow”?
  - All three words are spelled correctly
  - Enumerating each character: the space is large
  - Solution: using logs of queries, e.g., **flew from** vs. fled fore

Li et al. A generalized hidden Markov model with discriminative training for query spelling correction. SIGIR 2012

PageRank

- How to rank webpages?
  - Using retrieval models: only captures relevance

- Capturing quality of web pages:
  - Based on how often the page is cited
  - Intuition: a popular website (e.g., Google) would be cited by a lot of other webpages

PageRank

- ***"The Anatomy of a Large-Scale Hypertextual Web Search Engine"*** - *Sergey Brin and* Lawrence Page, *Computer networks and ISDN systems, 1998*

$$PR(p_i) = \frac{1-d}{N} + d \sum_{p_j \in M(p_i)} \frac{PR(p_j)}{L(p_j)}$$

- Favors pages that are highly cited, and pages cited by highly cited pages

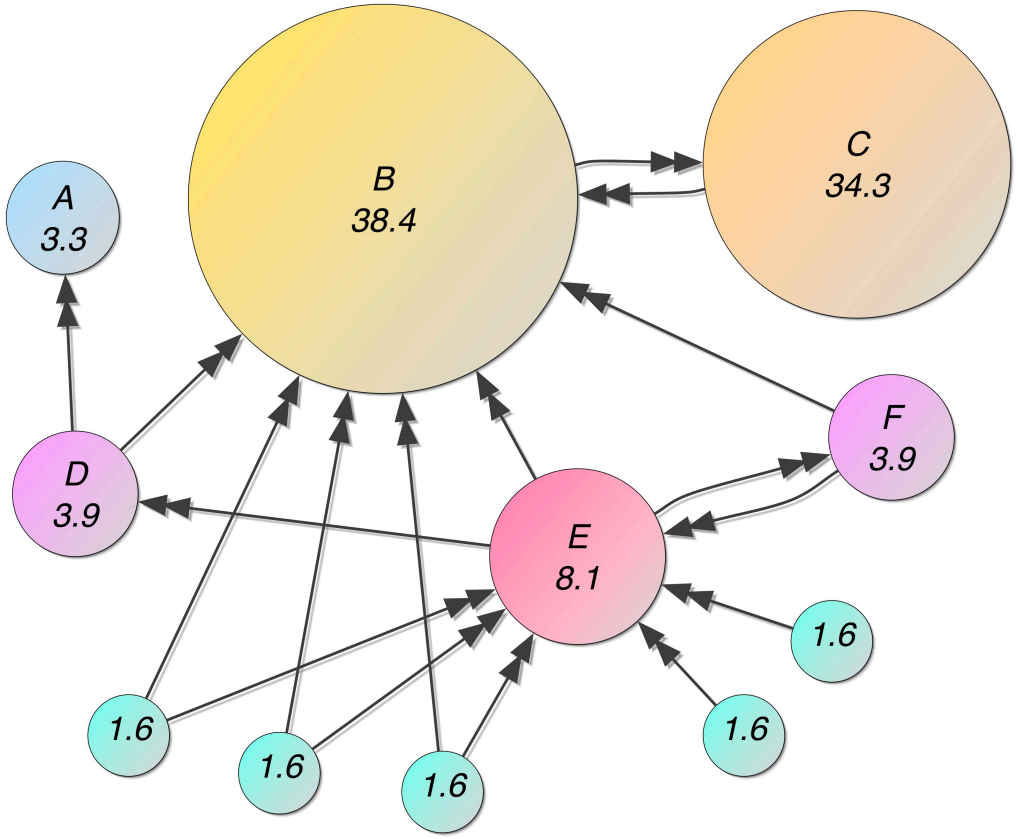

1/2 probability of randomly walking into B

PageRank

- Assign each node an initial page rank
- Repeat until convergence
  - Calculate the page rank of each node using the equation

$$PR(p_i) = \frac{1-d}{N} + d \sum_{p_j \in M(p_i)} \frac{PR(p_j)}{L(p_j)}$$

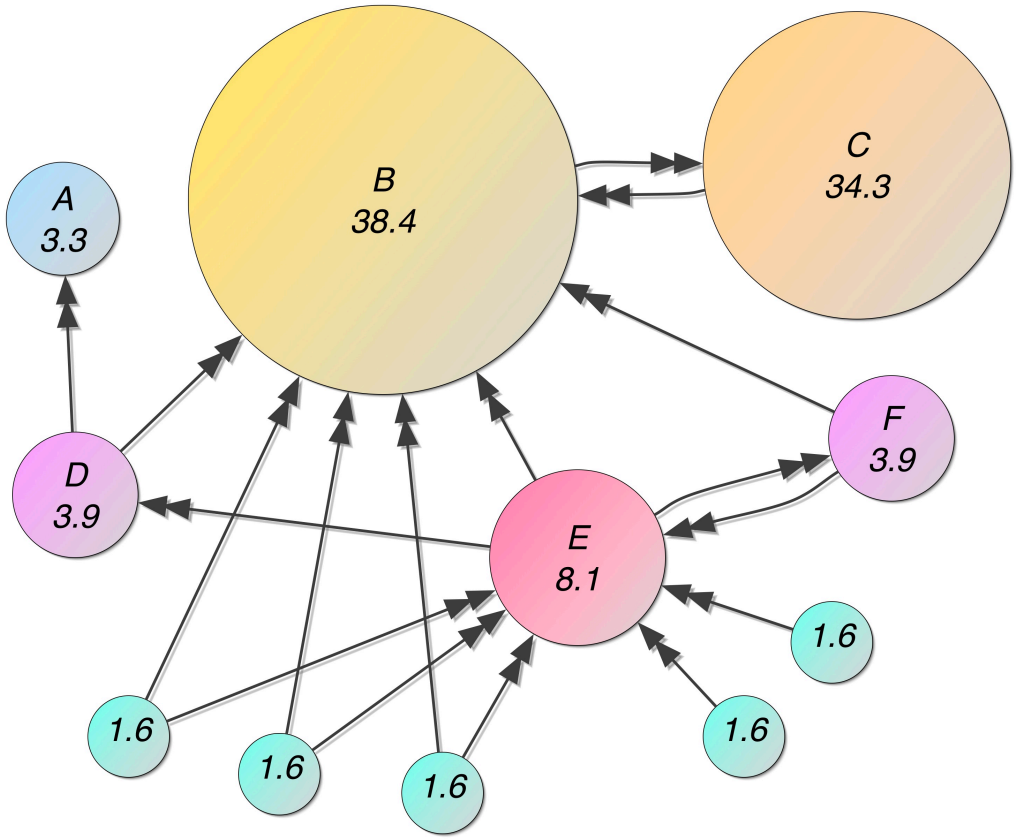

Problems of page rank

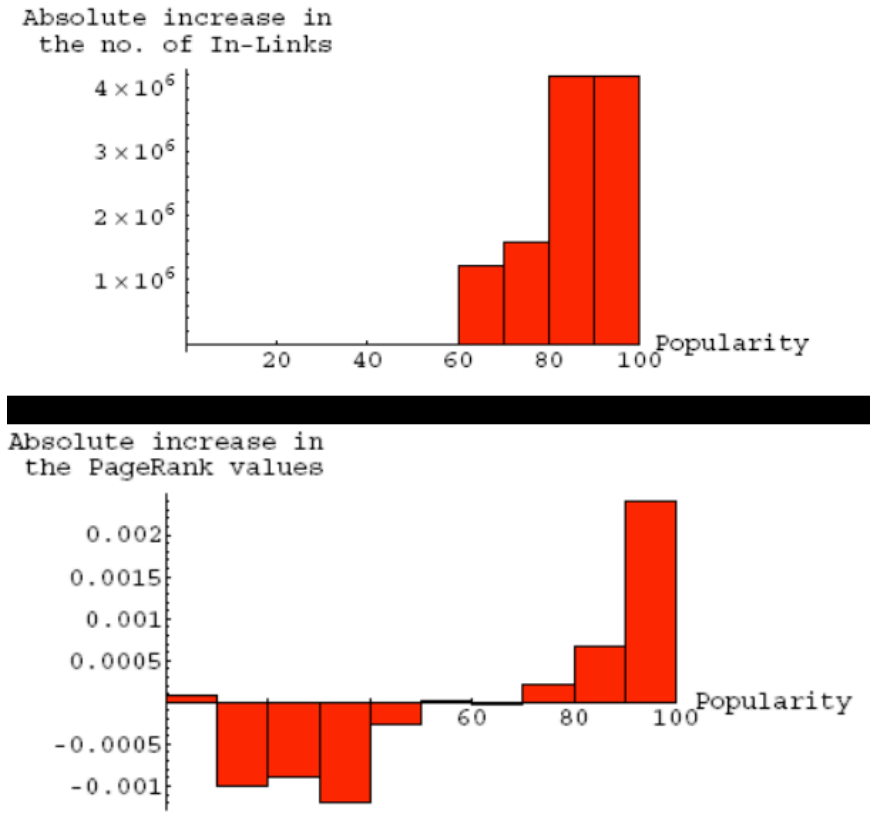

rich gets richer

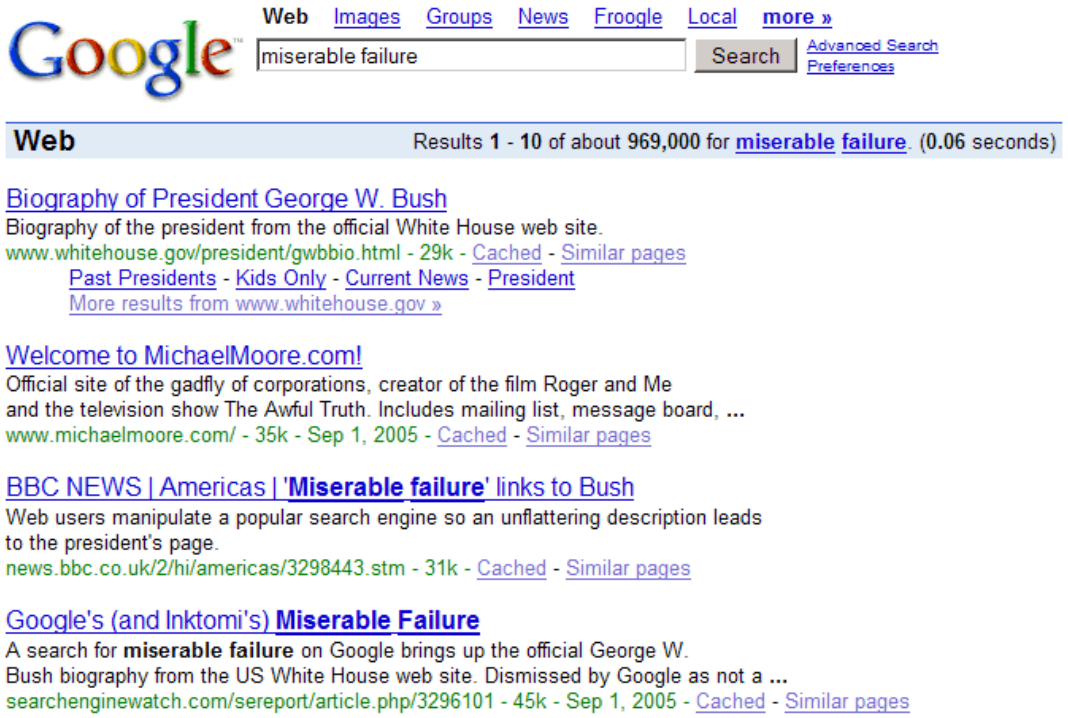

Google bombing

HITS

- Hubs: compilations of a broad catalog of information that led users direct to other authoritative pages

- Authorities: a page that is linked by many different hubs

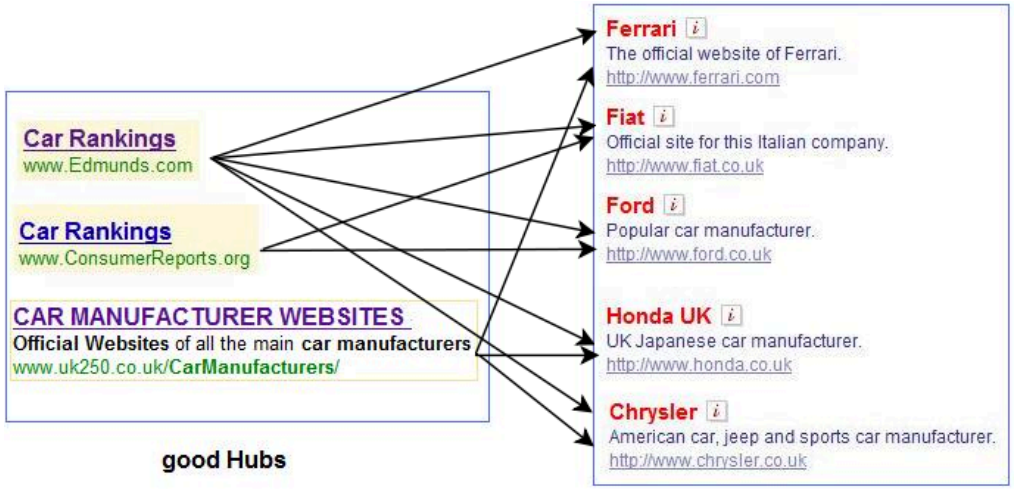

good Hubs

good Authorities

Query: Top automobile makers

HITS

- Repeat k times
  - Update hub score: v = A^T u
  - Update authority score: u = A^T v

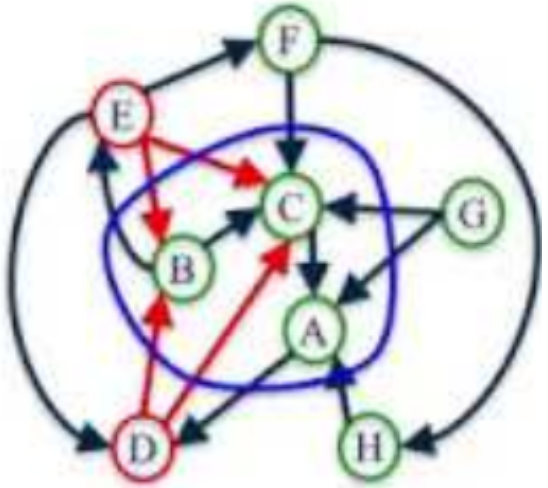

Search engine tools

- Apache Lucene
  - Free and open search engine library
  - First developed in 1999
- ElasticSearch
  - A search engine
  - based on Lucene

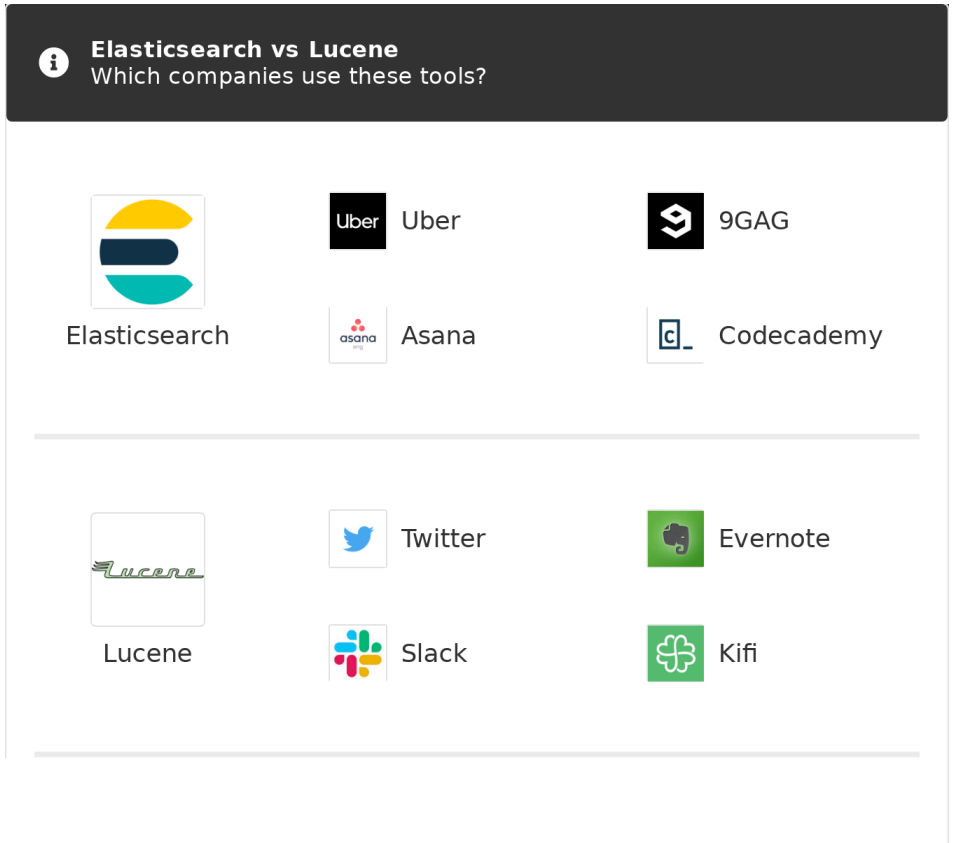

ElasticSearch

- Using a REST api

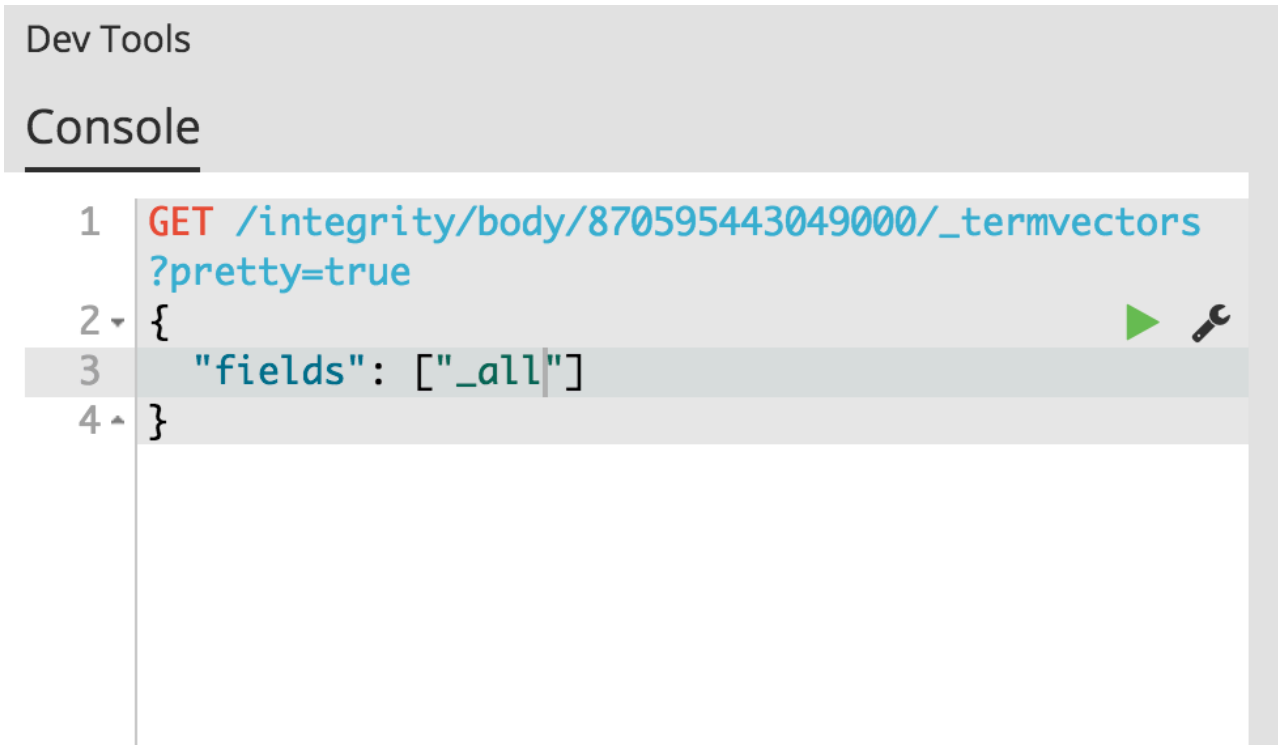

Homework 2: Using ElasticSearch to build a search engine

- Build an inverted index
- Evaluate three search algorithm's performance
  - TF-IDF
  - BM25
  - Dirichlet-LM